

AEL — Technical Blueprint

1. The Parts

The Context Layer: The segment profile — mid-sized, seasonal Dutch/Benelux wholesalers with high product-variant complexity, broken size runs, and club/league calendar-driven ordering cycles — and the unified quality criteria (Maturity Tiers, ROI anchors, security limitations) are contained in a static configuration file that loads first on each run.

Ingestion & Translation: Uses an LLM call to translate raw inputs (Dutch URLs, text files, or transcript text) into technical B1/B2 English and eliminates unnecessary market noise. vendor jurisdiction/ownership (EU vs. US/Chinese AI vendor, data-residency and CLOUD Act exposure) and internal AI-maturity signals (what a source's decision-makers already understand about AI-based forecasting, and whether there is a shared internal sense of what "good enough" historical data looks like). Only interview transcripts, not desk-research documents, can supply the internal AI-maturity signal, since it depends on what the company's own people say, not on published material.

Appraisal & Extraction: Only information pertinent to amateur teamwear operational KPIs (such as cut-to-ship cycles and fabric waste) is extracted by comparing the translated language to the Context Layer.

Drafting: Enforces the "minimum 2 text pages before a graphic" requirement by organizing the extracted insights into our standard Markdown handbook template.

Evaluation Pipeline: An automated process that examines the drafted Markdown to determine the Flesch-Kincaid readability score and confirm that the required "Quick-Start Action Sheet" is present.

Review & Publish: Before performing the git commit and git push to the GitHub Pages wiki, a human evaluates the final Markdown.

2. What passes between the parts

The strict contract between the Appraisal part and the Drafting part is a structured JSON file (e.g., `extracted_insights.json`). It always contains specific keys.

3. The boundary that nothing can pass

All unanonymized interview audio recordings from the site visit and proprietary corporate data (such as actual historical ERP sales numbers) must be locally anonymized before being uploaded to any public or external LLM API.

4. The PRD trace

PRD Promise: "autonomously handle translation from Dutch." -> Managed by the Ingestion & Translation section.

PRD Promise: "Output to our current wiki on GitHub." -> Managed by the Review & Publish section.

PRD Promise "Flesch-Kincaid reading level. Evaluated by the Evaluation Pipeline component.

PRD Promise: "Quantifiable ROI Anchors and Maturity Tiers." -> This is ensured by the Context Layer, strongly prompting the Appraisal portion.