

AEL — Knowledge Architecture

1. What the platform has to know

In order for our tool to write a manual for a sports equipment wholesaler, it needs to work with four basic “sets” of information:

- Permanent rules and context: A single master file that clearly describes who we’re writing the guide for (our SME profile) and what our quality rules are (e.g., the text must be readable and always include an action summary). This file doesn’t change much, and the agent must take it into account for every task.
- External online sources: Reports, case studies, and articles that we as a team find and download from the internet over the course of several weeks.
- Field data: Our own notes and anonymized transcripts from interviews we conduct directly at the company.
- Drafts: Texts generated by artificial intelligence that are waiting in the system for one of us to review and finally publish.

2. Where it lives, and how it is found again

Where we store it: We won’t be using any complex database systems. We’ll store all data as text files in well-organized folders (GitHub) directly within our project. This is the easiest to maintain, and anyone on the team can open and read the files at any time.

1. Ingestion - Local source folder

Reads raw Dutch whitepapers and Deventrade interview transcripts. A local script automatically removes staff names, club names, and private prices.

2. Appraisal - Prompt & context layer

Translates Dutch text into simple B1/B2 English. Extracts only facts matching amateur teamwear KPIs and saves them into `extracted_insights.json`.

3. Drafting - Markdown engine

Formats the insights into our handbook template (keeping text under 2 pages before a visual, plus the 1-page Action Sheet). Checks the readability score.

4. Publishing - CLI review & Git

A team member reviews the Markdown file manually. Once approved, we push it to the GitHub Pages wiki.

3. Why this, and not the alternatives

- **The alternative we considered:** We considered using a database.
- **Why we rejected it:** We decided to use ordinary files in folders because we promised in the PRD that anyone on the team would be able to use the tool. Any of us can open and read regular folders on a computer without any special knowledge
- **What would make us switch:** We would switch to a database only if, during the course, we collected such a huge amount of text that our tool could no longer easily filter it.

4. The line, kept

In our proposal, we have established a strict rule: we will never send actual company data or raw interview recordings to public artificial intelligence systems. These sensitive materials will therefore be stored in a completely separate folder. Our agent will be programmed so that, for security reasons, it will not be permitted to access this folder at all. Before any field data reaches the AI, they are processed by our local cleaning script so we don't have to do it manually.

The boundary therefore looks like this:

Protected	side:
------------------	--------------

Raw interview transcripts → staff names, club names, private prices

Local	anonymization	step:
--------------	----------------------	--------------

Sensitive information is removed or replaced before any model can access the text.

Model-readable	side:
-----------------------	--------------

Anonymized text → extracted insights → handbook draft

This setup helps us keep company information private while still allowing the platform to use useful knowledge from the interviews.

5. The open questions

We did not receive any questions or feedback from the Socratic tutor about our blueprint. Because of this, there were no specific changes required.